

# Realtime embedded coding under Linux

Bernd Porr

February 17, 2022

# Contents

<b>1</b>	<b>Writing C++ device driver classes</b>	<b>2</b>
1.1	Introduction	2
1.2	General recommendations how to write your C++ classes for devices	3
1.3	Low level userspace device access	4
1.3.1	SPI	4
1.3.2	I2C	5
1.3.3	Access GPIO pins	7
1.3.4	Access to hardware via special devices in /sys	8
1.3.5	I2S: Audio	9
1.3.6	Accessing physical memory locations (danger!)	10
1.4	Callbacks in C++ sensor classes	10
1.4.1	With pigpio library callback	10
1.4.2	Thread with a blocking “read” or “poll operation”	11
1.5	Kernel driver programming	12
1.6	Conclusion	13

# Chapter 1

## Writing C++ device driver classes




Figure 1.1: Dataflow and timing in low level realtime coding

### 1.1 Introduction

Fig. 1.1 shows the basic dataflow and how event timing is established. While it's obvious that data needs to flow from/to the hardware it's even more important to guarantee its timing. On the hardware-side the timing is guaranteed by event signals, data-ready signals and also by the timing of a serial interface. The Linux kernel translates this timing info into blocking I/O on pseudo filesystems such as /dev or /sys which means that a read operation blocks till data has arrived or an event has happened. Some low level libraries such as pigpio translate them back into callbacks but generally that needs to be done by you within a C++ class.

This lecture focusses on you writing your own C++ class hiding away the complexity (and messy) low level C APIs and raw device access to /dev and /sys.

## 1.2 General recommendations how to write your C++ classes for devices

As said above the main purpose of object oriented coding here is to hide away the complexity of low level driver access and offer the client a simple and safe way of connecting to the sensor. In particular:

1. Getters, setters and callbacks hand over *physical units* (temperature, acceleration, ...) and not raw integer values which have no meaning.
2. The sensor is configured by specifying physical units (time, voltage, temperature) and not sensor registers. Default values should be specified that it can start with little knowledge of the internals.
3. The class handles the realtime processing by offering callback interfaces to receive from and methods to transmit data to the sensor.
4. The class is re-usable outwith of your specific project and has its own cmake project (for example in a subdirectory), i.e. is a library. It has docstrings for all public functions and constants and has documentation generated by doxygen.

Keep SOLID in mind when writing your C++ device classes:

1. *Single responsibility*: If you have a temperature sensor and an accelerometer then write two classes, one for the temperature sensor and one of the accelerometer. In terms of
2. *Open-Closed principle and Dependency Inversion Principle*: Think what both sensors have in common and create a base class which inherits its functionality. For example if both sensors are SPI based you could have the basic SPI management in the base class and then the temperature sensor and accelerometer inherits it.
3. *Interface Segregation Principle*: Do not force a client to implement an abstract method and make it optional. Instead of adding an abstract callback to your sensor class itself rather add a method where the client can register and unregister the callback.

## 1.3 Low level userspace device access

The following sections provide pointers of how to write the C++ driver classes for different hardware protocols.

### 1.3.1 SPI

Table 1.1: SPI modes

SPI Mode	CPOL	CPHA	Idle state
0	0	0	L
1	0	1	L
2	1	1	H
3	1	0	H

Transfer to/from SPI is best managed by the low level access to `/dev`. Open the SPI device with the standard file operations:

```
fd = open(spiDevice, O_RDWR);
```

Then it's important to set the SPI mode (see table. 1.1):

```
ioctl(fd, SPI_IOC_WR_MODE, &mode);
```

which are explained, for example, here: <https://www.analog.com/en/analog-dialogue/articles/introduction-to-spi-interface.html>.

SPI transmits and receives at the same time so we need to use `ioctl` to do the communication. There is a special structure which needs to be filled:

```
struct spi_ioc_transfer tr = {
    .tx_buf = (unsigned long)tx1,
    .rx_buf = (unsigned long)rx1,
    .len = ARRAY_SIZE(tx1),
    .delay_usecs = delay,
    .speed_hz = speed,
    .bits_per_word = 8,
};
```

which points to two character buffers “tx” and “rx” with the same length.

Then open the SPI device in the /dev filesystem:

```
fd = open( "/dev/spidev0.0", O_RDWR);
```

Reading and simultaneous writing is happening then via the ioctl function:

```
ret = ioctl(fd, SPI_IOC_MESSAGE(1), &tr);  
if (ret < 1)  
    pabort("can't send spi message");
```

Sometimes the SPI protocol of a chip is so odd that even the raw I/O via /dev won't work and you need to write your own bit banging interface, for example done here for the ADC on the alphabot: <https://github.com/berndporr/alphabot/blob/main/alphabot.cpp#L58>. This is obviously far from ideal as it might require “usleep” commands so that acquisition needs to be run in a separate thread (the alphabot uses a timer callback in a separate thread).

Overall the SPI protocol is often device dependent and calls often for experimentation to get it to work. Often the SPI clock is also the ADC conversion clock which requires a longer lasting clock signal by transmitting dummy bytes in addition to the payload.

As a general recommendation do not use SAR converters which use the SPI data clock also as acquisition clock as they are often not compatible with the standard SPI transfers via /dev. Use sensors or ADCs which have their own clock signal.

### 1.3.2 I2C

The I2C bus has two signal lines (SDA & SDL) which must be pulled up by resistors. Every I2C device has an address on the bus. You can scan a bus with the “i2cdetect” (part of the i2c-tools package):

```
root@raspberrypi:/home/pi# i2cdetect -y 1  
    0  1  2  3  4  5  6  7  8  9  a  b  c  d  e  f  
00:  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  
10:  --  --  --  --  --  --  --  --  --  --  --  --  --  1e  --  
20:  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  
30:  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  
40:  --  --  --  --  --  --  --  --  --  --  --  --  --  --  --  
50:  --  --  --  --  --  --  --  --  58  --  --  --  --  --  --
```

```

60: -- -- -- -- -- -- -- -- -- -- 6b -- -- -- --
70: -- -- -- -- -- -- -- --
root@raspberrypi:/home/pi#

```

In this case there are 3 I2C devices on the bus at addresses 1E, 58 and 6B. The address needs to be specified when accessing the I2C device.

### Raw /dev/i2c access

I2C either transmits or receives but never at the same time so here we can use the standard C read/write commands. However, we need to use ioctl to tell the kernel to which I2C address we are talking:

```

char buf[2];
int file = open("/dev/i2c-2", O_RDWR);
int addr = 0x58;
ioctl(file, I2C_SLAVE, addr);
write(file, buf, 1)
read(file, buf, 2)

```

where “addr” is the I2C address. Then use standard read() or write() commands. Usually the 1st write() operation tells the chip which register to read or write to. Then write/read its register.

### I2C access via pigpio

Access via pigpio (<http://abyz.me.uk/rpi/pigpio/cif.html>) is preferred in contrast to the direct access of the raw /dev/i2c because many different devices can be connected to the I2C bus and pigpio manages this. Simply install the development package:

```
sudo apt-get install libpigpio-dev
```

which triggers then the install of the other relevant packages. For example writing a byte to a register in an I2C sensor can be done with a few commands:

```

int fd = i2cOpen(i2c_bus, address, 0);
i2cWriteByteData(fd, subAddress, data);
i2cClose(fd);

```

where i2c\_bus is the I2C bus number (usually one on the RPI) and the address is the I2C address of the device on that bus. The subAddress here is the register address in the device.

### 1.3.3 Access GPIO pins

#### **/sys filesystem**

The GPIO of the raspberry PI can easily be controlled via the /sys filesystem. This is slow but good for debugging as you can directly write a numerical “0” or “1” to it and print the result. The pseudo files are here:

`/sys/class/gpio`

`/sys` contains files which directly relate to individual pins. To be able to access a pin we need to tell Linux to make it visible:

`/sys/class/gpio/export`

For example, writing a 5 (in text form) to this file would create the subdirectory `/sys/class/gpio/gpio5` for GPIO pin 5.

Then reading from

`/sys/class/gpio/gpio5/value`

would give you the status of GPIO pin 5 and writing to it would change it.

**GPIO interrupt handling via /sys** The most important application for the /sys filesystem is to do interrupt processing in userspace. A thread can be put to sleep until an interrupt has happened on one of the GPIO pins. This is done by monitoring the “value” of a GPIO pin in the /sys filesystem with the “poll” command:

```
struct pollfd fdset[1];
int nfds = 1;
int gpio_fd = open("/sys/class/gpio/gpio5/value", O_RDONLY | O_NONBLOCK );
memset((void*)fdset, 0, sizeof(fdset));
fdset[0].fd = gpio_fd;
fdset[0].events = POLLPRI;
int rc = poll(fdset, nfds, timeout);
if (fdset[0].revents & POLLPRI) {
    // dummy
    read(fdset[0].fd, buf, MAX_BUF);
}
```

makes the program go to sleep until an interrupt has occurred on GPIO pin 5. Then the program wakes up and continues.



## pigpio

The above section gives you a deep understanding what's happening on the /sys - level but it's recommended to use the pigpio library (<http://abyz.me.uk/rpi/pigpio/cif.html>) to read/write to GPIO pins or do interrupt programming.

For example to set GPIO pin 24 as an input just call:

```
gpioSetMode(24,PI_INPUT);
```

To read from GPIO pin 24 just call:

```
int a = gpioRead(24)
```

**interrupt handling via pigpio** pigpio can call a function whenever a change has occurred on a GPIO pin. A function of the form: "static void gpioISR(int gpio, int level, uint32\_t tick, void\* userdata)" can be registered with pigpio:

```
gpioSetISRFuncEx(24,RISING_EDGE,ISR_TIMEOUT,gpioISR,(void*)this);
```

where "this" is the pointer to your class instance which is then used to call a class method, here: "dataReady()".

```
class LSM9DS1 {
    void dataReady();
    static void gpioISR(int gpio, int level, uint32_t tick, void* userdata)
    {
        if (level) {
            ((LSM9DS1*)userdata)->dataReady();
        }
    }
};
```

See [https://github.com/berndporr/LSM9DS1\\_RaspberryPi\\_CPP\\_Library](https://github.com/berndporr/LSM9DS1_RaspberryPi_CPP_Library) for the code.

### 1.3.4 Access to hardware via special devices in /sys

Some sensors are directly available via the sys filesystem in human readable format.

For example

```
cat /sys/class/thermal/thermal_zone0/temp
```

gives you the temperature of the CPU.

### 1.3.5 I2S: Audio

The standard framework is also: <https://github.com/alsa-project>.

ALSA works packet based where a read command returns a chunk of audio or a chunk is written to.

First the parameters are requested and the driver can modify or reject them:

```
/* Signed 16-bit little-endian format */
snd_pcm_hw_params_set_format(handle, params,
                             SND_PCM_FORMAT_S16_LE);

/* One channel (mono) */
snd_pcm_hw_params_set_channels(handle, params, 1);

/* 44100 bits/second sampling rate (CD quality) */
val = 44100;
snd_pcm_hw_params_set_rate_near(handle, params,
                                 &val, &dir);
```

Then playing sound is done in an endless loop where a read() or write() command is issued. Both are blocking so that it needs to run in a thread:

```
while(running) {
    rc = snd_pcm_writei(handle, buffer, frames);
    if (rc == -EPIPE) {
        /* EPIPE means underrun */
        fprintf(stderr, "underrun occurred\n");
        snd_pcm_prepare(handle);
    } else if (rc < 0) {
        fprintf(stderr,
                "error from writei: %s\n",
                snd_strerror(rc));
    } else if (rc != (int)frames) {
        fprintf(stderr,
                "short write, write %d frames\n", rc);
    }
}
```

For a full coding example “aplay” is a very good start or “arecord” which can be found here: <https://github.com/alsa-project>.

### 1.3.6 Accessing physical memory locations (danger!)

In case you really need to access registers you can access also memory directly. This should only be used as a last resort. For example, setting the clock for the AD converter requires turning a GPIO pin into a clock output pin. This is not yet supported by the drivers so we need to program registers on the RPI.

- Linux uses virtual addressed so that a pointer won't point to a physical address. It points to three page tables with an offset.
- Special device `/dev/mem` which allows access of physical memory.
- The command "mmap" provides a pointer to a physical address by opening `/dev/mem`.
- Example:

```
int *addr;
if ((fd = open("/dev/mem", O_RDWR|O_SYNC)) < 0 ) {
    printf("Error opening file. \n");
    close(fd);
    return (-1);
}
addr = (int *)mmap(0, num*STRUCT_PAGE_SIZE, PROT_READ, MAP_PRIVATE,
                  fd, 0x0000620000000000);
printf("addr: %p \n",addr);
printf("addr: %d \n",*addr);
```

- Use this with care! It's dangerous if not used properly.

## 1.4 Callbacks in C++ sensor classes

### 1.4.1 With pigpio library callback

As shown above pigpio provides a callback mechanism for GPIO interrupts so that you just need to offer that callback to the main program. That's done by a *callback* class or often called "interface":

```
class LSM9DS1callback {
public:
    virtual void hasSample(LSM9DS1Sample sample) = 0;
};
```

The main program then implements the abstract method “hasSample()”, instantiates the class and then saves its pointer in the driver class, here called “lsm9ds1Callback”. The pigpio callback then simply calls the method “hasSample()”:

```
void LSM9DS1::dataReady() {
    LSM9DS1Sample sample;
    // fills the sample struct with data
    // ...
    lsm9ds1Callback->hasSample(sample);
}
```

## 1.4.2 Thread with a blocking “read” or “poll operation”

Again we’ll be handing the data over via a *callback* class:

```
class AD7705callback {
public:
    virtual void hasSample(int sample) = 0;
};
```

and the header of the main class looks like:

```
class AD7705Comm {
    AD7705Comm(const char* spiDevice = "/dev/spidev0.0");
    void setCallback(AD7705callback* cb);
    void start();
    void stop();
}
```

where we provide a pointer to the callback class.

The main program then creates a derived class from the abstract callback class and then registers this:

```

class AD7705printSampleCallback : public AD7705callback {
    virtual void hasSample(int v) {
        // process the sample here
    }
};

```

Finally create a separate thread which runs in a loop where the thread is put to sleep either by “poll()” or “read()” and woken up periodically after a data ready signal as been received:

```

SysGPIO dataReadyGPIO(22);
int sysfs_fd = dataReadyGPIO.gpio_fd_open();
ad7705comm->running = 1;
while (ad7705comm->running) {
    dataReadyGPIO.gpio_poll(sysfs_fd,1000);
    int value = ad7705comm->readData(ad7705comm->fd)-0x8000;
    if (ad7705comm->ad7705callback) {
        ad7705comm->ad7705callback->hasSample(value); }
    }
close(sysfs_fd);

```

where the GPIO port 22 is monitored for new data by a thin wrapper class for the /sys filesystem found here: <https://github.com/berndporr/gpio-sysfs>. The flag “running” keeps the loop going till the main program stops it by setting it to false and then the thread terminates. The thread goes to sleep at “gpio\_poll” and wakes up when the data-ready pin triggers GPIO pin 22. Then, the data is read and if the callback has been registered by the main program “hasSample” is called.

The full example is here: [https://github.com/berndporr/rpi\\_AD7705\\_daq](https://github.com/berndporr/rpi_AD7705_daq).

## 1.5 Kernel driver programming

This is beyond the scope of this lecture but we can help you to write your own Linux driver if needed. This is for example the case if you do high speed data acquisition higher than 10kHz.

## 1.6 Conclusion

From the sections above it's clear that Linux userspace low level device access is complex, even without taking into account the complexity of contemporary sensors which have often a multitude of registers and pages of documentation. Your task is to hide away all this (scary) complexity in a C++ class and offer the client an easy to understand interface.