Realtime embedded coding under Linux

Bernd Porr & Nick Bailey

March 23, 2022

Contents

Chapter 1

Introduction

Figure 1.1: Dataflow and timing in low level realtime coding

In realtime embedded coding it's all about *events* which can be a binary signal as somebody opening a door or an ADC signalling that a sample is ready. Fig. 1.1 shows the basic dataflow and how event timing is established: devices by themselves have event signals such as data ready or crash sensor triggered. The Linux kernel then receives interrupt callbacks. However, userspace has no direct interrupt mechanism but it has blocking I/O where an I/O operation blocks until an interrupt has happened. These can then be translated back into callbacks between classes. Data is transmitted back to the hardware via methods called "setters".

Fig. 1.2 shows the overall communication between C++ classes in a realtime system. This communication is done via callbacks (*not* getters) and setters, where an event from a sensor traverses according to its realtime requirements through the classes via callbacks and then back to the control output via setters. For example, a collision sensor at a robot triggers a GPIO pin, which then triggers

Figure 1.2: A realtime system with two C++ classes. Communication between classes is achieved with callbacks (not getters) for incoming events and setters to send out control events. The control output itself receives its timing from the events so that the loop is traversed as quickly as possible.

a callback to issue an avoidance action which in turn then sets the motors in reverse.

When developing the C++ classes keep S.O.L.I.D. https://www.digitalocean. com/community/conceptual_articles/s-o-l-i-d-the-first-five-principles-of-object-o in mind:

1. *Single responsibility*: If you have a temperature sensor and an accelerometer then write two classes, one for the temperature sensor and one of the accelerometer. Don't write one class `Hardware` which takes an argument `Hardware::temp` or `Hardware::accel` and then does one of two completely different things. It's a debugging nightmare.

2. *Open-Closed principle* Your class is open to extension but closed to modification. For example an ADC class has a callback which returns voltage to the client. However, you'll be connecting, for example, a temperature sensor to it, so you'd like to be able to extend the class overloading the callback methods so that you add the conversion from volt to degrees. This is a Good Thing™ so long as you create a derived class. It's a bad idea to hack the existing ADC class adding a `to_kelvin` method: why would somebody using the ADC to read a value from an accelerometer need that?

3. *Liskov substitution principle* Strictly, subtituting a derived class for its base class does not result in the program becoming incorrect. That is to say, any derived class from your device driver class can be used in place of the base class if the base class is all that's required, because the extra functionality in the derived class shouldn't break the basic required functionality of the

base class. For example, if you have a super duper DAC with lots of extra features, it shouldn't stop you using it when you only need a very simple one. This also means that sensible default values should be set so that the client won't need to understand the nerdy features of that super duper DAC.

4. *Interface Segregation Principle*: Keep functionality separate and rather divide it up in different classes. Imagine you have a universal IO class with SPI and I2C but your client really just needs SPI. Then the client is forced to deactivate I2C or in the worst case the class causes collateral damage without the client knowing why.

5. *Dependency inversion*: That is about obstructing the essential features of of a class of interfaces. For example, ideally you want a base class covering a range of similar ADC converters from the same manufacturer and not a base class being a driver for a particular chip. Individual ADC chip driver classes then inherit from the abstract ADC driver. The authors of the ADC's driver will no doubt consider their chip's particular capabilities to be the core ideas, but this needs inverting. Why would one ADC driver need to provide all the necessaray code for a different one of the same family? Instead, the driver is dependent on the abstract idea of an ADC, not the other way around.

Besides S.O.L.I.D it's also essential that:

1. the project has a **build system** such as cmake. It's strongly recommended to use cmake (autoconf only for older existing projects).

2. classes (in particular driver classes) are **re-usable** outwith of the specific project and have their own "cmake" projects in their subdirectories. Except for the main cmake project all sub-projects are libraries.

3. all public interfaces have **doc-strings** for all public methods/constants and an automatically generated reference, for example with the help of doxygen.

4. classes which perform internal processing such as filters, databases, detectors, . . . have **unit tests** and run via "ctest".

5. the documentation provides comprehensive information about the project itself, how to install and run the project.

Chapter 2

Writing C++ device driver classes

This chapter focuses on writing your own C++ device driver class hiding away the complexity (and messy) low level C APIs and/or raw device access. How are events translated into I/O operations? On the hardware-side we have event signals such as data-ready signals or by the timing of a serial or audio interface. The Linux kernel translates this timing info into blocking I/O on pseudo filesystems such as /dev or /sys which means that a read operation blocks until data has arrived or an event has happened. Some low level libraries such as pigpio translate them back into C callbacks. The task of a C++ programmer is to hide this complexity and these quite different approaches in C++ classes which communicate via callbacks and setters with the client classes.

2.1 General recommendations how to write your C++ classes for devices

As said above the main purpose of object oriented coding here is to hide away the complexity of low level driver access and offer the client a simple and safe way of connecting to the sensor. In particular:

1. Setters and callbacks hand over *physical units* (temperature, acceleration, . . .) or relative units but not raw integer values with no meaning.

2. The sensor is configured by specifying physical units (time, voltage, temperature) and not sensor registers. Default config parameters should be specified that the class can be used straight away with default parameters.

3. It has simple demo programs demonstrating how to use the class by a client.

2.2 Low level userspace device access

The following sections provide pointers of how to write the C++ driver classes for different hardware protocols.

2.2.1 SPI

Table 2.1: SPI modes

SPI Mode	CPOL	CPHA	Idle state
0	0	0	L
1	0	1	L
2	1	1	H
3	1	0	H

SPI is a protocol which usually transmits and receives at the same time. Even that data might not be used it needs to be matched up. So for example sending 8 byes and receiving 8 bytes at the same time.

Transfer to/from SPI is best managed by the low level access to /dev. Open the SPI device with the standard open() function:

```
int fd = open( "/dev/spidev0.0", O_RDWR);
```

Then set the SPI mode (see table. 2.1):

```
int ret = ioctl(fd, SPI_IOC_WR_MODE, &mode);
```

which is explained, for example, here: https://www.analog.com/en/analog-dialogue/articles/introduction-to-spi-interface.html.

Since SPI transmits and receives at the same time we need to use ioctl to do the communication. Populate this struct:

```
struct spi_ioc_transfer tr = {
  .tx_buf = (unsigned long)tx1,
```

```
    .rx_buf = (unsigned long)rx1,
    .len = ARRAY_SIZE(tx1),
    .delay_usecs = delay,
    .speed_hz = speed,
    .bits_per_word = 8,
};
```

which points to two character buffers "tx" and "rx" with the same length.

Reading and simultaneous writing is happening then via the ioctrl function:

```
int ret = ioctl(fd, SPI_IOC_MESSAGE(1), &tr);
```

Sometimes the SPI protocol of a chip is so odd that even the raw I/O via /dev won't work and you need to write your own bit banging interface, for example done here for the ADC on the alphabot: https://github.com/berndporr/alphabot/blob/main/alphabot.cpp#L58. This is obviously far from ideal as it might require "usleep" commands so that acquisition needs to be run in a separate thread (the alphabot uses a timer callback in a separate thread).

Overall the SPI protocol is often device dependent and calls for experimentation to get it to work. Often the SPI clock is also the ADC conversion clock which requires a longer lasting clock signal by transmitting dummy bytes in addition to the payload.

As a general recommendation do not use SAR converters which use the SPI data clock also as acquisition clock as they are often not compatible with the standard SPI transfers via /dev. Use sensors or ADCs which have their own clock signal.

2.2.2 I2C

The I2C bus has two signal lines (SDA & SDL) which must be pulled up by resistors. Every I2C device has an address on the bus. You can scan a bus with "i2cdetect" (part of the i2c-tools package):

```
root@raspberrypi:/home/pi# i2cdetect -y 1
     0  1  2  3  4  5  6  7  8  9  a  b  c  d  e  f
00:                         -- -- -- -- -- -- -- --
10: -- -- -- -- -- -- -- -- -- -- -- -- -- -- 1e --
20: -- -- -- -- -- -- -- -- -- -- -- -- -- -- -- --
30: -- -- -- -- -- -- -- -- -- -- -- -- -- -- -- --
```

```
40: -- -- -- -- -- -- -- -- -- -- -- -- -- -- -- --
50: -- -- -- -- -- -- -- -- 58 -- -- -- -- -- -- --
60: -- -- -- -- -- -- -- -- -- -- -- 6b -- -- -- --
70: -- -- -- -- -- -- -- --
root@raspberrypi:/home/pi#
```

In this case there are 3 I2C devices on the I2C bus at addresses 1E, 58 and 6B and need to be specified when accessing the I2C device.

Raw /dev/i2c access

I2C either transmits or receives but never at the same time so here we can use the standard C read/write commands. However, we need to use ioctrl to tell the kernel the I2C address:

```
char buf[2];
int file = open("/dev/i2c-2",O_RDWR);
int addr = 0x58;
ioctl(file, I2C_SLAVE, addr);
write(file,buf,1)
read(file,buf,2)
```

where "addr" is the I2C address. Then use standard read() or write() commands. Usually the 1st write() operation tells the chip which register to read or write to. Then write/read its register.

I2C access via pigpio

Access via pigpio (http://abyz.me.uk/rpi/pigpio/cif.html) is preferred in contrast to direct access of the raw /dev/i2c because many different devices can be connected to the I2C bus and pigpio manages this. Simply install the development package:

```
sudo apt-get install libpigpio-dev
```

which triggers then the install of the other relevant packages. For example writing a byte to a register in an I2C sensor can be done with a few commands:

```
int fd = i2cOpen(i2c_bus, address, 0);
i2cWriteByteData(fd, subAddress, data);
i2cClose(fd);
```

where i2c_bus is the I2C bus number (usually 1 on the RPI) and the address is the I2C address of the device on that bus. The subAddress here is the register address in the device.

2.2.3 Access GPIO pins

/sys filesystem

The GPIO of the raspberry PI can easily be controlled via the /sys filesystem. This is slow but good for debugging as you can directly write a "0" or "1" string to it and print the result. The pseudo files are here:

```
/sys/class/gpio
```

which contains files which directly relate to individual pins. To be able to access a pin we need to tell Linux to make it visible:

```
/sys/class/gpio/export
```

For example, writing a 5 (in text form) to this file would create the subdirectory /sys/class/gpio/gpio5 for GPIO pin 5.

Then reading from

```
/sys/class/gpio/gpio5/value
```

would give you the status of GPIO pin 5 and writing to it would change it. A thin wrapper around the GPIO sys filesystem is here: https://github.com/berndporr/gpio-sysfs.

GPIO interrupt handling via /sys   The most important application for the /sys filesystem is to do interrupt processing in userspace. A thread can be put to sleep until an interrupt has happened on one of the GPIO pins. This is done by monitoring the "value" of a GPIO pin in the /sys filesystem with the "poll" command:

```
struct pollfd fdset[1];
int nfds = 1;
int gpio_fd = open("/sys/class/gpio/gpio5/value", O_RDONLY | O_NONBLOCK );
memset((void*)fdset, 0, sizeof(fdset));
fdset[0].fd = gpio_fd;
```

```
fdset[0].events = POLLPRI;
int rc = poll(fdset, nfds, timeout);
if (fdset[0].revents & POLLPRI) {
    // dummy
    read(fdset[0].fd, buf, MAX_BUF);
}
```

makes the thread go to sleep until an interrupt has occurred on GPIO pin 5. Then the thread wakes up and execution continues.

pigpio

The above section has given you a deep understanding what's happening under the hood on the sysfs-level but it's highly recommended to use the pigpio library (http://abyz.me.uk/rpi/pigpio/cif.html) to read/write to GPIO pins or do interrupt programming.

For example to set GPIO pin 24 as an input just call:

```
gpioSetMode(24,PI_INPUT);
```

To read from GPIO pin 24 just call:

```
int a = gpioRead(24)
```

interrupt handling via pigpio   pigpio manages GPIO interrupt handling by wrapping all the above functionality into a single command where the client registers a callback function which is called whenever a change has occurred on a GPIO pin. Specifically a method of the form:

```
class mySensorClass {
  ...
  static void gpioISR(int gpio, int level, uint32\_t tick, void* userdata)
  ...
}
```

is registered with pigpio:

```
gpioSetISRFuncEx(24,RISING_EDGE,ISR_TIMEOUT,gpioISR,(void*)this);
```

where "this" is the pointer to your class instance which is then used to call a class method, here: "dataReady()".

```
class LSM9DS1 {
  void dataReady();
  static void gpioISR(int gpio, int level, uint32_t tick, void* userdata)
    {
        ((LSM9DS1*)userdata)->dataReady();
    }
};
```

where here within the static function the void pointer is cast back into the instance pointer. See https://github.com/berndporr/LSM9DS1_RaspberryPi_CPP_Library for the complete code.

2.2.4   Access to hardware via special devices in /sys

Some sensors are directly available via the sys filesystem in human readable format.

For example

```
cat /sys/class/thermal/thermal_zone0/temp
```

gives you the temperature of the CPU.

2.2.5   I2S: Audio

The standard framework for audio is alsa: https://github.com/alsa-project.

ALSA works packet based where a read command returns a chunk of audio or a chunk is written to.

First, the parameters are requested and the driver can modify or reject them:

```
/* Signed 16-bit little-endian format */
  snd_pcm_hw_params_set_format(handle, params,
                                 SND_PCM_FORMAT_S16_LE);

  /* One channel (mono) */
  snd_pcm_hw_params_set_channels(handle, params, 1);

  /* 44100 bits/second sampling rate (CD quality) */
  val = 44100;
  snd_pcm_hw_params_set_rate_near(handle, params,
                                    &val, &dir);
```

Then playing sound is done in an endless loop were a read() or write() command is issued. Both are blocking so that it needs to run in a thread:

```
while(running) {
    rc = snd_pcm_writei(handle, buffer, frames);
    if (rc == -EPIPE) {
      /* EPIPE means underrun */
      fprintf(stderr, "underrun occurred\n");
      snd_pcm_prepare(handle);
    } else if (rc < 0) {
      fprintf(stderr,
              "error from writei: %s\n",
              snd_strerror(rc));
    }  else if (rc != (int)frames) {
      fprintf(stderr,
              "short write, write %d frames\n", rc);
    }
}
```

For a full coding example "aplay" is a very good start or "arecord". Both can be found here: https://github.com/alsa-project.

2.2.6   Accessing physical memory locations (danger!)

In case you really need to access registers you can access also memory directly. This should only be used as a last resort. For example, setting the clock for the AD converter requires turning a GPIO pin into a clock output pin. This is not yet supported by the drivers so we need to program registers on the RPI.

- Linux uses virtual addressed so that a pointer won't point to a physical address. It points to three page tables with an offset.

- Special device /dev/mem which allows access of physical memory.

- The command "mmap" provides a pointer to a physical address by opening /dev/mem.

- Example:

```
int *addr;
if ((fd = open("/dev/mem", O_RDWR|O_SYNC)) < 0 ) {
    printf("Error opening file. \n");
    close(fd);
    return (-1);
}
addr = (int *)mmap(0, num*STRUCT_PAGE_SIZE, PROT_READ, MAP_PRIVATE,
            fd, 0x0000620000000000);
printf("addr: %p \n",addr);
printf("addr: %d \n",*addr);
```

- Use this with care! It's dangerous if not used properly.

2.3   Kernel driver programming

You can also create your own /dev/mydevice in the /dev filesystem by writing
a kernel driver and a matching userspace library. For example the USB mouse
has a driver in kernel space and translates the raw data from the mouse into
coordinates. However, this is beyond the scope of this handout. If you want to
embark on this adventure then the best approach is to find a kernel driver which
does approximately what you want and then modify it for your purposes.

2.4   Callbacks in C++ device classes

As said in the introduction your hardware device class has callback interfaces to
hand back the data to the client.

There are different ways of tackling the issue of callbacks but the simplest
one is defining a method as *abstract* and asking the client to implement it in a
derived class. That abstract function can either be in a separate interface class
or part of the device class itself. So, we have two options:

1. The callback is part of the device driver class:

```
class MyDriver {
        void start(DevSettings settings = DevSettings() );
        void stop();
        virtual void callback(float sample) = 0;
};
```

2. The callback is part of an interface class:

```
class CallbackInterface {
        virtual void callback(float sample) = 0;
};
```

and then registering it in the main device driver class:

```
class MyDriver {
        void registerCallback(CallbackInterface* cb);
};
```

These two options are now explained in greater detail.

2.4.1 Creating a callback interface

Here, we create a separate interface class containing a callback as an abstract method:

```
class LSM9DS1callback {
public:
        virtual void hasSample(LSM9DS1Sample sample) = 0;
};
```

The client then implements the abstract method "hasSample()", instantiates the interface class and then saves its pointer in the device class, here called "lsm9ds1Callback".

```
void LSM9DS1::dataReady() {
        LSM9DS1Sample sample;
        // fills the sample struct with data
        // ...
        lsm9ds1Callback->hasSample(sample);
}
```

The pointer to the interface instance is transmitted via a setter which receives the pointer of the interface as an argument, for example:

```
        void registerCallback(LSM9DS1callback* cb);
```

This allows to register a callback optionally. The client might not need one or not always. See https://github.com/berndporr/rpi_AD7705_daq for a complete example.

2.4.2 Adding directly an abstract method to the device driver class

Instead of creating a separate class containing the callback you can also add the callback straight to the device driver class.

```
class ADS1115rpi {
        ...
        virtual void hasSample(float sample) = 0;
        ...
};
```

This forces the client to implement the callback to be able to use the class. This creates a very safe environment as all dependencies are set at compile time and the abstract nature of the base class makes clear what needs to be implemented. See https://github.com/berndporr/rpi_ads1115 for a complete example.

2.4.3 Callback arguments

Above the callbacks just delivered one floating point value. However, often more than one sample or more complex data is transmitted:

- Complex data: Do not put loads of arguments into the callback but define a *struct*. For example an ADC might deliver all 4 channels at once:

```
class ADmulti {

        struct ADCSample {
            float ch1;
            float ch2;
            float ch3;
            float ch4;
        };


        ...
        virtual void hasSample(ADCSample sample) = 0;
        ...
};
```

- Arrays: Use arrays which contain the length of the arrays: either std::array, std::vector, etc or const arrays and then references to these so that the callback knows the length. For example the LIDAR callback uses a reference to a const length array:

```
/**
 * Callback interface which needs to be implemented by the user.
 **/
struct DataInterface {
        virtual void newScanAvail(
                float rpm,
                A1LidarData (&)[A1Lidar::nDistance]) = 0;
};
```

  where "A1Lidar::nDistance" defines a reference to a constant length array which in itself consists of "A1LidarData" stucts.

In terms of *memory management*:

1. Low sampling rate complex data structures: allocate as a local variable. It can be a simple type or a struct. See "dataReady()" in: https://github.com/berndporr/LSM9DS1_RaspberryPi_CPP_Library/blob/master/LSM9DS1.cpp.

2. High sampling rate buffers: allocate memory on the heap in the constructor or in the private section of the class as a const length array and pass on a *reference*. See "getData()" in https://github.com/berndporr/rplidar_rpi.

2.5   Conclusion

In conclusion, the communication between C++ is done via callbacks and setters where the event from the sensor traverses through the C++ classes via callbacks and then back to the control output via setters.

From the sections above it's clear that Linux userspace low level device access is complex, even without taking into account the complexity of contemporary chips which have often a multitude of registers and pages of documentation. Your task is to hide away all this (scary) complexity in a C++ class and offer the client an easy to understand interface.

Chapter 3

Threads

3.1 Introduction

In a realtime system events need to be dealt with as soon as possible while also interacting with the user and performing other background tasks. However, certain operations might take considerable computing time (FFT, etc) or I/O is blocking while waiting for data. The solution are threads.

3.2 Processes and Threads

Processes are different programs which seem to be running at the same time. However this is done by the operating system which switches approximately every 10ms from one process to the next so it feels as if they are running at the same time. A thread is a lightweight process where multiple threads share the same memory and started from within the parent process. As with processes the threads seem to be running at the same time. When a thread is started it runs simultaneously to the main process which created it.

3.3 Thread and worker

A thread is just a *container* for the actual method which is running independently. The method contained inside of a thread is often called *worker*.

3.3.1 Creating threads

In C++ a worker is a method within a class and needs to be *static* which means it won't be able to access the instance variables of a class. The trick is to pass a pointer to the instance of the class ("this") as the argument of the worker, for example, called "exec":

```
uthread = new std::thread(MyClassWithAThread::exec, this);
```

where `MyClassWithAThread` is a class containing the static function "exec":

```
class MyClassWithAThread {
        void run() {
                // ... hard work is done here
                doCallback(result); // hand the result over
        }
        static void exec(MyClassWithAThread* cppThread) {
                cppThread->run();
        }
}
```

which in turn then calls a non-static class method "run()" which then has access to the instances variables.

3.3.2 Lifetime of a thread

Threads terminate simply once the static worker has finished its job. To tell the client that a thread has finished you can use a *callback* to trigger an event.

Sometimes it's important to wait for the termination of the thread, for example when your whole program is terminating or when you stop an endless loop in a thread. To wait for the termination of the thread use the "join()" method:

```
        void stop() {
                uthread->join();
                delete uthread;
        }
```

Important is also to release the memory of a thread after it has finished to avoid memory leaks.

3.3.3 Running/stopping workers with endless loops

Threads with endless loops are often used in conjunction with blocking I/O which provide the timing:

```
void run() {
      running = true;
      while (running) {
             read(buffer); // blocking
             doCallback(buffer); // hand data to client
      }
}
```

Note the flag "running" which is controlled by the main program and is set to zero to terminate the thread:

```
      void stop() {
             running = false; // <----- HERE!!
             uthread->join();
             delete uthread;
      }
```

Note that "join()" is a blocking operation and needs to be used with care not to lock up the main program. You probably only need it when your program is terminating. See https://github.com/berndporr/rpi_AD7705_daq for an example.

3.3.4 Timing within threads

Threads are perfect to create timing without using sleep commands with the help of *blocking I/O*.

select/poll commands waiting for GPIO interrupts

In section 2.2.3 we introduced the so called "poll" command which is not polling an IRQ pin but *putting a thread to sleep* till an external event has happened. Then of course a callback function should be called reacting to the external event. This is the preferred method for low latency responses.

As said previously, use *pigpio* on the Raspberry PI which wraps the select/poll commands into a thread and calls a *callback* function whenever an GPIO pin has been triggered.

Timing with blocking I/O

Blocking I/O (read, write, etc) can be used to time the data coming in because the thread goes to sleep when it's waiting for I/O but wakes up very quickly after new data has arrived.

In this example the blocking "read" command creates the timing of the callback:

```
void run() {
        running = 1;
        while (running) {
                read(buffer); // blocking
                doCallback(buffer); // hand data to client
        }
}
```

Timing with Linux/pigpio timers

Similar to threads one can create timers which are called at certain intervals. These timers emit a Linux signal at a specified interval and then this signal is caught by a global (static) function. Generally it's *not recommended* to use timers for anything which needs to be reliably sampled, for example ADC converters or sensors with sampling rates higher than a few Hz. On the raspberry PI use the pigpio library and its timer callbacks – if needed at all.

Chapter 4

Realtime/event processing in QT

4.1 Introduction

QT is a cross platform windows development environment for Linux/Windows
and Mac.

Elements in QT are *Widgets* which can contain anything form plots, buttons
or text fields. They are classes. You can define your own widgets or use ready
made ones.

4.2 Layout in QT

There are different ways of declaring layout in QT. One is using a markup lan-
guage which then has matching classes or creating it all with classes. We show
how it works by just using classes which organise the layout. There is also a trend
in general to use code to declare the layout as done in SwiftUI, for example.

This is an example how widgets are organised into nested vertical and hori-
zontal layouts (see Fig. 4.1 for the result).

```
// create 3 widgets
button = new QPushButton;
thermo = new QwtThermo;
plot = new QwtPlot;

// vertical layout
vLayout = new QVBoxLayout;
```

Figure 4.1: QT example layout

```
vLayout->addWidget(button);
vLayout->addWidget(thermo);

// horizontal layout
hLayout = new QHBoxLayout;
hLayout->addLayout(vLayout);
hLayout->addWidget(plot);

// main layout
setLayout(hLayout);
```

4.3 Callbacks in QT

4.3.1 Events from widgets

In contrast to our low level callback mechanism using interfaces QT rather directly calls methods in classes. The problem is that function pointers cannot be directly used as a class has instance pointers to its local data. So a method of a class needs to be combined with the instance pointer. The QT method "connect" does exactly that:

```
connect(button,&QPushButton::clicked,
        this,&Window::reset);
```

The QPushButton instance "button" has a method called "clicked" which is called whenever the user clicks on the button. This is then forwarded to the method "reset" in the application Widget.

4.3.2   Plotting realtime data arriving via a callback

The general idea is to store the realtime samples from a callback in a buffer and trigger a screen refresh at a lower rate for example every 40 ms which then plots the contents of the buffer.

The callback "addSample" here is called in realtime whenever a sample has arrived:

```
void Window::addSample( float v ) {
    // add the new input to the plot
    memmove( yData, yData+1, (plotDataSize-1) * sizeof(double) );
    yData[plotDataSize-1] = v;
}
```

which stores the sample 'v' in the shift buffer 'yData'.

Then the screen refresh (which is slow) is done at a lower and unreliable rate:

```
void Window::timerEvent( QTimerEvent * )
{
    curve->setSamples(xData, yData, plotDataSize);
    plot->replot();
    thermo->setValue( yData[0] );
    update();
}
```

After "update()" has been called in the timer event the paint event is executed by QT as soon as possible and re-paints the canvas of the widget:

```
void ScopeWindow::paintEvent(QPaintEvent *) {
        QPainter paint( this );

        paint.drawLine( ... )
}
```

Note that neither the timer nor the "update()" function is called in a reliable way but whenever QT likes to do it. So the QT timers cannot be used to sample data but should only be used for screen refresh.

Chapter 5

Realtime web server/client communication

5.1 Introduction

Web server / client applications are at the heart of a huge number of web applications ranging from shopping baskets to social media applications.

Generally it's easy to create dynamic content with many different solutions available (see PHP or nodejs) and well documented. However, feeding realtime data from C++ to a web page or realtime button presses back to C++ is a bit more difficult.

Important to recognise where *events* are generated: it's always the client (web browser, mobile app) which triggers an event, be it sending data over to the server or requesting data. It's always initiated by the client.

5.2 REST

The interface between a web client (browser, phone app) is usually implemented as a Representational State Transfer Architectural (REST) style API by communicating via an URL on a web server. The most important requirements for this API are very general and won't define the actual data format:

1. **Uniform interface** Any device connecting to the URL should get the same reply. No matter if a web page or mobile phone requests the temperature of a sensor the returned format must always be the same.

2. **Client-server decoupling** The only information the client needs to know is the URL of the server to request data or send data.

3. **Statelessness** Each request needs to include all the information necessary and must not depend on previous requests. For example a request to a buffer must not alter the buffer but just read from it so that another user reading the buffer shortly after receives the same data.

See https://www.ibm.com/cloud/learn/rest-apis for the complete list of REST design principles.

5.3 Data formats

5.3.1 Server → client: JSON (application/json)

The most popular dataformat is JSON which is basically a map of key/value pairs which can also be nested:

```
{
    temperature: [20, 21, 20, 19, 17],
    steps: 100,
    comment: "all good!"
}
```

Since JSON is sort of readable text a web server can simply generate that text send it over via http or https. There is no difference except that the MIME format is now 'application/json' instead of html.

5.3.2 Client → server: POST (application/x-www-form-urlencoded)

When a website sends data back to the server it needs to encode it in the form of a single text-line where the key/value pairs are combined with &-signs:

```
temperature=20&steps=100&comment=all+good%33
```

The receiver then has the task to entangle this stream into a suitable dataformat, for example a map. All server side scripting languages such as PHP or node.js have powerful functions to decode these strings. In C/C++ libcurl can be used for decoding.

5.4 Server

On the Linux system a web server needs to be set up. There are a variety of different options available but we are focusing here on the ones which can be used for C++ communication (i.e. CGI).

5.4.1 Web servers (http/https)

- NGINX: Easy to configure but very flexible web server

- Apache: Hard to configure but safe option

- lighttpd: Smaller web server with a small memory footprint

Note that it's possible to run different web servers at the same time where they then act as proxies for a central web server visible to the outside world. In particular nginx makes it very easy to achieve this.


Figure 5.1: FastCGI dataflow.

5.4.2 FastCGI

FastCGI (see Fig 5.1) is written in C++ and generates the entire content of the http/https request. In particular here we generate JSON packets server side which can then be processed by client JavaScripts. For realtime applications JSON transmission is perfect because the client-side JavaScript can request JSON packages and directly turn them into variables.

A fast CGI program is a UNIX commandline program which communicates with the web server (nginx, Apache, . . .) via a UNIX socket which in turn is a pseudo file located in a temporary directory for example '/tmp/sensorsocket'.

The web server then maps certain http/https requests to his socket, for example, the configuration for nginx looks like this:

```
location /sensor/ {
    include        fastcgi_params;
    fastcgi_pass   unix:/tmp/sensorsocket;
  }
```

If the user does a request via the URL `www.mywebpage.com/sensor/` then nginx contacts the fastcgi program via this socket. The fastcgi program then needs to return the content. Internally this will be a C++ callback inside of the fastcgi program.

The C++ fastcgi API https://github.com/berndporr/json_fastcgi_web_api is wrapper around the quite cryptic fastcgi C library and we are discussing its callback handlers now.

Server → client: JSON (application/json)

The fastCGI callback expects a JSON string with the data transmitted form the server to the client. There is helper class JSONGenerator which generates the JSON data from various C++ types:

```
class JSONcallback : public JSONCGIHandler::GETCallback {
public:
/**
* Gets the data and sends it to the webserver.
* The callback creates two JSON entries. One with the
* timestamp and one with the temperature from the sensor.
**/
virtual std::string getJSONString() {
    JSONCGIHandler::JSONGenerator jsonGenerator;
    jsonGenerator.add("epoch",(long)time(NULL));
    jsonGenerator.add("temperatures",temperatureArray);
    return jsonGenerator.getJSON();
    }
};
```

5.4.3   Client → server: POST (application/x-www-form-urlencoded)

Like in any GUI the client can press a button and create an event. On the client side this is packaged as a JSON record with jquery and then sent over to the

server. The server then receives the data as a callback:

```
virtual void postString(std::string postArg) {
auto m = JSONCGIHandler::postDecoder(postArg);
float temp = atof(m["volt"].c_str());
std::cerr << m["hello"] << "\n";
sensorfastcgi->forceValue(temp);
}
```

5.5   Client

Generally on the client side (= web page) HTML with embedded *JavaScript* is used to generate realtime output/input without reloading the web page. JavaScript is *event driven* and has callbacks so it's perfect for realtime applications. Use jQuery to request and post JSON from/to the server.

For example here we request data from the server as a JSON packet every second:

```
// callback when the JSON data has arrived
function getterCallback(result) {
  var temperatureArray = result.temperatures;
  // plot the array here
}

// timer callback (same idea as in QT to define a refresh rate)
function getTemperature() {
  // get the JSON data
  $.getJSON("/data/:80",getterCallback);
}

// document ready callback
function documentReady() {
  // request new data from the server every second
  window.intervalId = setInterval(getTemperature , 1000);
}

// called when the web page has been loaded
$(document).ready( documentReady );
```

Chapter 6

Setters

In Fig. 1.2 we have seen that data flows from the sensors to the C++ classes via *callbacks* then it flows back from the inner C++ classes to motor or display outputs is via *setters*. Setters are also used for setting configuration parameters.

A setter is a simple method in a class, for example to set the speed of a motor:

```
class Motor {
  /**
  * Set the Left Wheel Speed
  * @param speed between -1 and +1
  **/
  void setLeftWheelSpeed(float speed);
};
```

Again as with callbacks it's important to *abstract* away from the hardware, for example normalising the speed of the motor between -1 and $+1$ and *hiding* away the complexity of the PWM or GPIO ports in the class.

If a setter has more than one argument, in particular for configuration, it's highly recommended to use a *struct* to set the values. For example setting the parameters of the ADS1115:

```
/**
 * ADS1115 initial settings when starting the device.
 **/
struct ADS1115settings {
```

```
/**
 * I2C bus used (99% always set to one)
 **/
int i2c_bus = 1;

/**
 * I2C address of the ads1115
 **/
uint8_t address = DEFAULT_ADS1115_ADDRESS;
};

/**
 * Starts the data acquisition in the background and the
 * callback is called with new samples.
 * \param settings A struct with the settings.
 **/
void start(ADS1115settings settings = ADS1115settings() );
```

If a setter sets large buffers then it's highly recommended to allocate the memory in the constructor of the class and then call the setter by reference while running. Use array types which convey their length, for example std::array or a standard const array which implicitly carries their length.