

# Realtime embedded coding under Linux

Bernd Porr

March 2, 2022

# Contents

<b>1</b>	<b>Writing C++ device driver classes</b>	<b>2</b>
1.1	Introduction	2
1.2	General recommendations how to write your C++ classes for devices	3
1.3	Low level userspace device access	5
1.3.1	SPI	5
1.3.2	I2C	6
1.3.3	Access GPIO pins	8
1.3.4	Access to hardware via special devices in /sys	10
1.3.5	I2S: Audio	10
1.3.6	Accessing physical memory locations (danger!)	11
1.4	Kernel driver programming	12
1.5	Callbacks in C++ device classes	12
1.5.1	Creating a callback interface	13
1.5.2	Adding directly an abstract method to the device driver class	14
1.6	Conclusion	14
<b>2</b>	<b>Threads</b>	<b>15</b>
2.1	Introduction	15
2.2	Processes and Threads	15
2.2.1	Creating threads	15
2.2.2	Lifetime of a thread	16
2.2.3	Running/stopping threads with endless loops	16
2.2.4	Dealing with competition / concurrency	17
2.2.5	Timing within threads	17

# Chapter 1

## Writing C++ device driver classes




Figure 1.1: Dataflow and timing in low level realtime coding

### 1.1 Introduction

Fig. 1.1 shows the basic dataflow and how event timing is established. While it's obvious that data needs to flow from/to the hardware it's even more important to guarantee its timing in realtime applications. On the hardware-side the timing is guaranteed by event signals, data-ready signals and also by the timing of a serial interface. The Linux kernel translates this timing info into blocking I/O on pseudo filesystems such as /dev or /sys which means that a read operation blocks till data has arrived or an event has happened. Some low level libraries such as pigpio translate them back into C callbacks. Your task is to hide this complexity and these quite different approaches in C++ classes which communicate via callbacks and setters with the client classes.




Figure 1.2: A realtime system with two C++ classes. Communication between classes is achieved with callbacks (not getters) for incoming events and setters to send out control events. The control output itself receives its timing from the events so that the loop is traversed as quickly as possible.

Fig. 1.2 shows the overall communication between C++ classes in a realtime system. This communication is done via callbacks (*not* getters) and setters where an event from a sensor traverses according to its realtime requirements through the classes via callbacks and then back to the control output via setters. For example a collision sensor at a robot triggers a GPIO pin, which then triggers a callback to issue an avoidance action which in turn then sets the motors in reverse.

This chapter focuses on writing your own C++ device driver class hiding away the complexity (and messy) low level C APIs and/or raw device access to /dev and /sys.

## 1.2 General recommendations how to write your C++ classes for devices

As said above the main purpose of object oriented coding here is to hide away the complexity of low level driver access and offer the client a simple and safe way of connecting to the sensor. In particular:

1. Setters and callbacks hand over *physical units* (temperature, acceleration, ...) and not raw integer values which have no meaning.
2. The sensor is configured by specifying physical units (time, voltage, temperature) and not sensor registers. Default config parameters should be specified that the class can be used straight away with default parameters.

3. The class handles the realtime processing by offering callback interfaces (i.e. based on classes with virtual and abstract methods) to transmit data from the sensor and setters to transmit data to the device.
4. A build system is essential. It's strongly recommended to use "cmake" (autoconf only for older existing projects).
5. The class is re-usable outwith of your specific project and has its own "cmake" project in a subdirectory and is a library.
6. It has doc-strings for all public functions and constants, has documentation generated by doxygen.
7. It has simple demo programs demonstrating how to use the class by a client.

Keep S.O.L.I.D. [https://www.digitalocean.com/community/conceptual\\_articles/s-o-l-i-d-the-first-five-principles-of-object-oriented-design](https://www.digitalocean.com/community/conceptual_articles/s-o-l-i-d-the-first-five-principles-of-object-oriented-design) in mind when writing your C++ device classes:

1. *Single responsibility*: If you have a temperature sensor and an accelerometer then write two classes, one for the temperature sensor and one of the accelerometer. In terms of
2. *Open-Closed principle* Your class is open to extension but closed to modification. For example an ADC class has a callback which returns voltage to the client. However, you'll be connecting for example a temperature sensor to it so you'd like to be able to extend the class by for example overloading the callback methods so that you add the conversion from volt to degrees but not hacking the existing ADC class.
3. *Liskov substitution principle* Any derived class from your device driver class can be used in place of the base class if the base class is all that's required, because the extra functionality in the derived class shouldn't break the basic required functionality of the base class. For example, if you have a super duper DAC with lots of extra features, it shouldn't stop you using it when you only need a very simple one. This also means that sensible default values should be set so that the client won't need to understand the nerdy features of that super duper DAC.

4. *Interface Segregation Principle*: Keep functionality separate and rather divide it up in different classes. Imagine you have a universal IO class with SPI and I2C but your client really just needs SPI. Then the client is forced to deactivate I2C or in the worst case the class causes collateral damage without the client knowing why.
5. *Dependency inversion*: That is about obstructing the essential features of a class of interfaces. For example ideally you want, for example, a base class for covering a range of similar ADC converters from the same manufacturer and not a base class being a driver for a particular chip. Instead you want all ADC chip driver classes to inherit from the abstract ADC driver.

## 1.3 Low level userspace device access

The following sections provide pointers of how to write the C++ driver classes for different hardware protocols.

### 1.3.1 SPI

Table 1.1: SPI modes

SPI Mode	CPOL	CPHA	Idle state
0	0	0	L
1	0	1	L
2	1	1	H
3	1	0	H

SPI is a protocol which usually transmits and receives at the same time. Even that data might not be used it needs to be matched up. So for example sending 8 bytes and receiving 8 bytes at the same time.

Transfer to/from SPI is best managed by the low level access to /dev. Open the SPI device with the standard `open()` function:

```
int fd = open( "/dev/spidev0.0", O_RDWR);
```

Then set the SPI mode (see table. [1.1](#)):

```
int ret = ioctl(fd, SPI_IOC_WR_MODE, &mode);
```

which is explained, for example, here: <https://www.analog.com/en/analog-dialogue/articles/introduction-to-spi-interface.html>.

Since SPI transmits and receives at the same time we need to use `ioctl` to do the communication. Populate this struct:

```
struct spi_ioc_transfer tr = {  
    .tx_buf = (unsigned long)tx1,  
    .rx_buf = (unsigned long)rx1,  
    .len = ARRAY_SIZE(tx1),  
    .delay_usecs = delay,  
    .speed_hz = speed,  
    .bits_per_word = 8,  
};
```

which points to two character buffers “tx” and “rx” with the same length.

Reading and simultaneous writing is happening then via the `ioctl` function:

```
int ret = ioctl(fd, SPI_IOC_MESSAGE(1), &tr);
```

Sometimes the SPI protocol of a chip is so odd that even the raw I/O via `/dev` won't work and you need to write your own bit banging interface, for example done here for the ADC on the alphabot: <https://github.com/berndporr/alphabot/blob/main/alphabot.cpp#L58>. This is obviously far from ideal as it might require “usleep” commands so that acquisition needs to be run in a separate thread (the alphabot uses a timer callback in a separate thread).

Overall the SPI protocol is often device dependent and calls for experimentation to get it to work. Often the SPI clock is also the ADC conversion clock which requires a longer lasting clock signal by transmitting dummy bytes in addition to the payload.

As a general recommendation do not use SAR converters which use the SPI data clock also as acquisition clock as they are often not compatible with the standard SPI transfers via `/dev`. Use sensors or ADCs which have their own clock signal.

## 1.3.2 I2C

The I2C bus has two signal lines (SDA & SDL) which must be pulled up by resistors. Every I2C device has an address on the bus. You can scan a bus with “i2cdetect” (part of the `i2c-tools` package):

```

root@raspberrypi:/home/pi# i2cdetect -y 1
    0  1  2  3  4  5  6  7  8  9  a  b  c  d  e  f
00:                -- -- -- -- -- --
10: -- -- -- -- -- -- -- -- -- -- -- 1e --
20: -- -- -- -- -- -- -- -- -- -- -- -- --
30: -- -- -- -- -- -- -- -- -- -- -- -- --
40: -- -- -- -- -- -- -- -- -- -- -- -- --
50: -- -- -- -- -- -- -- -- 58 -- -- -- -- --
60: -- -- -- -- -- -- -- -- -- -- 6b -- -- --
70: -- -- -- -- -- -- -- -- -- -- -- -- --
root@raspberrypi:/home/pi#

```

In this case there are 3 I2C devices on the I2C bus at addresses 1E, 58 and 6B and need to be specified when accessing the I2C device.

### Raw /dev/i2c access

I2C either transmits or receives but never at the same time so here we can use the standard C read/write commands. However, we need to use ioctl to tell the kernel the I2C address:

```

char buf[2];
int file = open("/dev/i2c-2",O_RDWR);
int addr = 0x58;
ioctl(file, I2C_SLAVE, addr);
write(file,buf,1)
read(file,buf,2)

```

where "addr" is the I2C address. Then use standard read() or write() commands. Usually the 1st write() operation tells the chip which register to read or write to. Then write/read its register.

### I2C access via pigpio

Access via pigpio (<http://abyz.me.uk/rpi/pigpio/cif.html>) is preferred in contrast to direct access of the raw /dev/i2c because many different devices can be connected to the I2C bus and pigpio manages this. Simply install the development package:

```

sudo apt-get install libpigpio-dev

```



which triggers then the install of the other relevant packages. For example writing a byte to a register in an I2C sensor can be done with a few commands:

```
int fd = i2cOpen(i2c_bus, address, 0);
i2cWriteByteData(fd, subAddress, data);
i2cClose(fd);
```

where `i2c_bus` is the I2C bus number (usually 1 on the RPI) and the `address` is the I2C address of the device on that bus. The `subAddress` here is the register address in the device.

### 1.3.3 Access GPIO pins

#### **/sys filesystem**

The GPIO of the raspberry PI can easily be controlled via the `/sys` filesystem. This is slow but good for debugging as you can directly write a “0” or “1” string to it and print the result. The pseudo files are here:

```
/sys/class/gpio
```

which contains files which directly relate to individual pins. To be able to access a pin we need to tell Linux to make it visible:

```
/sys/class/gpio/export
```

For example, writing a 5 (in text form) to this file would create the subdirectory `/sys/class/gpio/gpio5` for GPIO pin 5.

Then reading from

```
/sys/class/gpio/gpio5/value
```

would give you the status of GPIO pin 5 and writing to it would change it. A thin wrapper around the GPIO `sys` filesystem is here: <https://github.com/berndporr/gpio-sysfs>.

**GPIO interrupt handling via `/sys`** The most important application for the `/sys` filesystem is to do interrupt processing in userspace. A thread can be put to sleep until an interrupt has happened on one of the GPIO pins. This is done by monitoring the “value” of a GPIO pin in the `/sys` filesystem with the “poll” command:

```

struct pollfd fdset[1];
int nfds = 1;
int gpio_fd = open("/sys/class/gpio/gpio5/value", O_RDONLY | O_NONBLOCK );
memset((void*)fdset, 0, sizeof(fdset));
fdset[0].fd = gpio_fd;
fdset[0].events = POLLPRI;
int rc = poll(fdset, nfds, timeout);
if (fdset[0].revents & POLLPRI) {
    // dummy
    read(fdset[0].fd, buf, MAX_BUF);
}

```

makes the thread go to sleep until an interrupt has occurred on GPIO pin 5. Then the thread wakes up and execution continues.

## **pigpio**

The above section has given you a deep understanding what's happening under the hood on the sysfs-level but it's highly recommended to use the pigpio library (<http://abyz.me.uk/rpi/pigpio/cif.html>) to read/write to GPIO pins or do interrupt programming.

For example to set GPIO pin 24 as an input just call:

```
gpioSetMode(24,PI_INPUT);
```

To read from GPIO pin 24 just call:

```
int a = gpioRead(24)
```

**interrupt handling via pigpio** pigpio manages GPIO interrupt handling by wrapping all the above functionality into a single command where the client registers a callback function which is called whenever a change has occurred on a GPIO pin. Specifically a method of the form:

```

class mySensorClass {
    ...
    static void gpioISR(int gpio, int level, uint32_t tick, void* userdata)
    ...
}

```

is registered with pigpio:

```
gpioSetISRFuncEx(24,RISING_EDGE,ISR_TIMEOUT,gpioISR,(void*)this);
```

where “this” is the pointer to your class instance which is then used to call a class method, here: “dataReady()”.

```
class LSM9DS1 {
    void dataReady();
    static void gpioISR(int gpio, int level, uint32_t tick, void* userdata)
    {
        ((LSM9DS1*)userdata)->dataReady();
    }
};
```

where here within the static function the void pointer is cast back into the instance pointer. See [https://github.com/berndporr/LSM9DS1\\_RaspberryPi\\_CPP\\_Library](https://github.com/berndporr/LSM9DS1_RaspberryPi_CPP_Library) for the complete code.

### 1.3.4 Access to hardware via special devices in /sys

Some sensors are directly available via the sys filesystem in human readable format.

For example

```
cat /sys/class/thermal/thermal_zone0/temp
```

gives you the temperature of the CPU.

### 1.3.5 I2S: Audio

The standard framework for audio is alsa: <https://github.com/alsa-project>.

ALSA works packet based where a read command returns a chunk of audio or a chunk is written to.

First, the parameters are requested and the driver can modify or reject them:

```
/* Signed 16-bit little-endian format */
snd_pcm_hw_params_set_format(handle, params,
                              SND_PCM_FORMAT_S16_LE);
```

```

/* One channel (mono) */
snd_pcm_hw_params_set_channels(handle, params, 1);

/* 44100 bits/second sampling rate (CD quality) */
val = 44100;
snd_pcm_hw_params_set_rate_near(handle, params,
                                &val, &dir);

```

Then playing sound is done in an endless loop were a `read()` or `write()` command is issued. Both are blocking so that it needs to run in a thread:

```

while(running) {
    rc = snd_pcm_writei(handle, buffer, frames);
    if (rc == -EPIPE) {
        /* EPIPE means underrun */
        fprintf(stderr, "underrun occurred\n");
        snd_pcm_prepare(handle);
    } else if (rc < 0) {
        fprintf(stderr,
                "error from writei: %s\n",
                snd_strerror(rc));
    } else if (rc != (int)frames) {
        fprintf(stderr,
                "short write, write %d frames\n", rc);
    }
}

```

For a full coding example “`aplay`” is a very good start or “`arecord`”. Both can be found here: <https://github.com/alsa-project>.

### 1.3.6 Accessing physical memory locations (danger!)

In case you really need to access registers you can access also memory directly. This should only be used as a last resort. For example, setting the clock for the AD converter requires turning a GPIO pin into a clock output pin. This is not yet supported by the drivers so we need to program registers on the RPI.

- Linux uses virtual addressed so that a pointer won’t point to a physical address. It points to three page tables with an offset.

- Special device `/dev/mem` which allows access of physical memory.
- The command “`mmap`” provides a pointer to a physical address by opening `/dev/mem`.
- Example:

```
int *addr;
if ((fd = open("/dev/mem", O_RDWR|O_SYNC)) < 0 ) {
    printf("Error opening file. \n");
    close(fd);
    return (-1);
}
addr = (int *)mmap(0, num*STRUCT_PAGE_SIZE, PROT_READ, MAP_PRIVATE,
                  fd, 0x0000620000000000);
printf("addr: %p \n",addr);
printf("addr: %d \n",*addr);
```

- Use this with care! It's dangerous if not used properly.

## 1.4 Kernel driver programming

You can also create your own `/dev/mydevice` in the `/dev` filesystem by writing a kernel driver and a matching userspace library. For example the USB mouse has a driver in kernel space and translates the raw data from the mouse into coordinates. However, this is beyond the scope of this handout. If you want to embark on this adventure then the best approach is to find a kernel driver which does approximately what you want and then modify it for your purposes.

## 1.5 Callbacks in C++ device classes

As said in the introduction your hardware device class has callback interfaces to hand back the data to the client.

There are different ways of tackling the issue of callbacks but the simplest one is defining a method as *abstract* and asking the client to implement it in a derived class. That abstract function can either be in a separate interface class or part of the device class itself. So, we have two options:

1. The callback is part of the device driver class:

```
class MyDriver {
    void start(DevSettings settings = DevSettings() );
    void stop();
    virtual void callback(float sample) = 0;
};
```

2. The callback is part of an interface class:

```
class CallbackInterface {
    virtual void callback(float sample) = 0;
};
```

and then registering it in the main device driver class:

```
class MyDriver {
    void registerCallback(CallbackInterface* cb);
};
```

These two options are now explained in greater detail.

### 1.5.1 Creating a callback interface

Here, we create a separate interface class containing a callback as an abstract method:

```
class LSM9DS1callback {
public:
    virtual void hasSample(LSM9DS1Sample sample) = 0;
};
```

The client then implements the abstract method “hasSample()”, instantiates the interface class and then saves its pointer in the device class, here called “lsm9ds1Callback”.

```
void LSM9DS1::dataReady() {
    LSM9DS1Sample sample;
    // fills the sample struct with data
    // ...
    lsm9ds1Callback->hasSample(sample);
}
```

The pointer to the interface instance is transmitted via a setter which receives the pointer of the interface as an argument, for example:

```
void registerCallback(LSM9DS1callback* cb);
```

This allows to register a callback optionally. The client might not need one or not always. See [https://github.com/berndporr/rpi\\_AD7705\\_daq](https://github.com/berndporr/rpi_AD7705_daq) for a complete example.

### 1.5.2 Adding directly an abstract method to the device driver class

Instead of creating a separate class containing the callback you can also add the callback straight to the device driver class.

```
class ADS1115rpi {  
    ...  
    virtual void hasSample(float sample) = 0;  
    ...  
};
```

This forces the client to implement the callback to be able to use the class. This creates a very safe environment as all dependencies are set at compile time and the abstract nature of the base class makes clear what needs to be implemented. See [https://github.com/berndporr/rpi\\_ads1115](https://github.com/berndporr/rpi_ads1115) for a complete example.

## 1.6 Conclusion

In conclusion, the communication between C++ is done via callbacks and setters where the event from the sensor traverses through the C++ classes via callbacks and then back to the control output via setters.

From the sections above it's clear that Linux userspace low level device access is complex, even without taking into account the complexity of contemporary chips which have often a multitude of registers and pages of documentation. Your task is to hide away all this (scary) complexity in a C++ class and offer the client an easy to understand interface.

# Chapter 2

## Threads

### 2.1 Introduction

In a realtime system events need to be dealt with as soon as possible while also interacting with the user and performing other background tasks. However, certain operations might take considerable computing time (FFT, etc) or I/O might be blocking while waiting for data. The solution are threads.

### 2.2 Processes and Threads

Processes are different programs which seem to be running at the same time. However this is done by the operating system which switches approximately every 10ms from one process to the next so it feels as if they are running at the same time. A thread is a lightweight process where multiple threads share the same memory and started from within the parent process. As with processes the threads seem to be running at the same time. When a thread is started it runs simultaneously to the main process which created it.

#### 2.2.1 Creating threads

In C++ a thread is a method within a class which needs to be a *static* which means it won't be able to directly access the instance variables of a class. The trick is to pass a pointer to the instance of the class as the argument of the static method running the thread, for example called "exec" here:

```
uthread = new std::thread(MyClassWithAThread::exec, this);
```



where `MyClassWithAThread` is a class containing the static function “exec”:

```
class MyClassWithAThread {
    void run();
    static void exec(MyClassWithAThread* cppThread) {
        cppThread->run();
    }
}
```

which in turn then calls a non-static class method “run()” which then has access to the instances variables. See [https://github.com/berndporr/rpi\\_AD7705\\_daq](https://github.com/berndporr/rpi_AD7705_daq) for an example.

## 2.2.2 Lifetime of a thread

Threads terminate simply once the static thread method has finished. To tell the client that a thread has finished use another *callback* to trigger an event.

Sometimes it's important to wait for the termination of the thread, for example when your whole program is terminating or for example when you stop the data acquisition of the ADC. To wait for the termination of the thread use the “join()” method:

```
void stop() {
    uthread->join();
    delete uthread;
    uthread = NULL;
}
```

## 2.2.3 Running/stopping threads with endless loops

Threads with endless loops are often used in conjunction with blocking I/O which provide the timing:

```
void run() {
    running = 1;
    while (running) {
        read(buffer); // blocking
        doCallback(buffer); // hand data to client
    }
}
```

You need to implement a flag (here: “running”) which is set to zero to terminate the thread. This is then done a method which terminates the thread by setting “running” to zero before calling “join()”:

```
void stop() {
    running = 0; // <----- HERE!!
    uthread->join();
    delete uthread;
    uthread = NULL;
}
```

Generally a thread should terminate as soon as possible once the flag “running” has been set to zero.

## 2.2.4 Dealing with competition / concurrency

To avoid that two threads manipulating data at the same time one can employ a “mutex”:

```
std::mutex cmdMtx;
```

Thread 1:

```
cmdMtx.lock();
a = 1
cmdMtx.unlock();
```

Thread 2:

```
cmdMtx.lock();
if (a == 2) { do dangerous stuff };
cmdMtx.unlock();
```

## 2.2.5 Timing within threads

### select/poll commands

These commands put the thread to sleep until an interrupt has been triggered at a GPIO pin. This is the fastest way of reacting to an external event. ADC converters can trigger an interrupt whenever data is available. This is the preferred method for low latency and can be even tuned more by requesting a Linux realtime scheduler.

```

struct pollfd fdset[1];
int nfd = 1;
int rc;
char buf[MAX_BUF];

memset((void*)fdset, 0, sizeof(fdset));

fdset[0].fd = gpiofd;
fdset[0].events = POLLPRI;

rc = poll(fdset, nfd, timeout);

if (fdset[0].revents & POLLPRI) {
    // dummy read
    read(fdset[0].fd, buf, MAX_BUF);
}

```

...where `gpiofd` is a file descriptor pointing to a GPIO pin which is usually under `/sys/bus/gpio/`. This fragment waits till there has been an interrupt on an IO pin. In a similar fashion you can also monitor other file descriptors if they support the poll command (BTW: this is *not polling*).

Please use `pigpio` on a Raspberry PI which provides a wrapper for the code above and offers a simple callback interface.

## Blocking I/O

I/O operations can be blocking (thread goes to sleep and wakes up when data is available) or non-blocking where the I/O operation returns immediately (perhaps with an error). Blocking I/O can be used to time the data coming in as the thread goes to sleep when it's waiting for I/O but wakes up very quickly after new data has arrived.

In this example from above the blocking “read” command creates the timing of the callback.

```

void run() {
    running = 1;
    while (running) {
        read(buffer); // blocking
        doCallback(buffer); // hand data to client
    }
}

```

```
    }  
}
```

### **Timing with Linux/pigpio timers**

Similar to threads one can create timers which are called at certain intervals. These timers emit a Linux signal at a specified interval and then this signal is caught by a global (static) function. On the raspberry PI it's recommended to use the pigpio library and its timer callbacks which create the most reliable timing. However, anything beyond 50Hz approximately is most likely not precise.